

Aaro7 - Helping MSMEs Rise On Time

Job Title: Cloud Engineer – AWS Infrastructure (Fintech Platform)

Screening Test

Section 1: EC2 + App Hosting

1. (Hands-on) Spin up an EC2 instance and deploy a simple FastAPI app with one route (/ping → {"message": "pong"})

Host it using Gunicorn and Nginx
Ensure it is accessible on `http://your-ec2-ip/ping`

Attach screenshots of:

Your Nginx config
Running process (Gunicorn)
Browser output of the endpoint

Step 1 : Launch EC2 Instance

Step 2 : Connect to EC2

```
sudo yum update -y && sudo yum install -y  
python3-pip python3-venv nginx
```

Step 3 : Create FastAPI app

```
mkdir fastapi-app && cd fastapi-app  
python3 -m venv venv  
source venv/bin/activate  
pip install fastapi uvicorn gunicorn
```

Step 4 : Create main.py

```
from fastapi import FastAPI

app = FastAPI()

@app.get("/ping")
def ping():
    return {"message": "pong"}
```

Step 5 : Run Gunicorn

```
gunicorn main:app -k uvicorn.workers.UvicornWorker --bind 127.0.0.1:8000
```

Step 5 : Install I& Configure Nginx

```
sudo apt install nginx -y
```

```
sudo vi /etc/nginx/sites-available/fastapi
```

```
server {
    listen 80;
    server_name _;

    location / {
        proxy_pass http://127.0.0.1:8000;
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
    }
}
```

```
sudo amazon-linux-extras enable nginx1
```

```
sudo yum install nginx -y
```

```
sudo systemctl enable nginx
```

```
sudo systemctl start nginx
```

```
sudo nginx -t
```

```
nginx: configuration file /etc/nginx/nginx.conf test is successful
```

```
sudo vi /etc/nginx/nginx.conf
```

```
sudo systemctl restart nginx
```

OUTPUT :

Let Check through : <http://52.66.235.166/ping>

```
{  
  "message": "pong"  
}
```


Running process (Gunicorn)

```
[ec2-user@ip-172-31-3-152 ~]$ sudo -i  
[root@ip-172-31-3-152 ~]# ls  
[root@ip-172-31-3-152 ~]# cd /  
[root@ip-172-31-3-152 /]# ls  
bin boot dev etc fastapi-app home lib lib64 local media mnt opt proc root run sbin srv sys test_env usr var  
[root@ip-172-31-3-152 /]# ps aux | grep gunicorn  
root      2126  0.0  0.0 119424  948 pts/0    S+   18:42   0:00 grep --color=auto gunicorn  
root      3439  0.0  2.5 239156 25176 ?        S    12:04   0:02 /fastapi-app/venv/bin/python3 /fastapi-app/venv/bin/gunicorn main:app -k uvicorn.workers.UvicornWorker --bind 127.0.0.1:8000  
root      3442  0.1  4.4 347592 43316 ?        S1   12:04   0:43 /fastapi-app/venv/bin/python3 /fastapi-app/venv/bin/gunicorn main:app -k uvicorn.workers.UvicornWorker --bind 127.0.0.1:8000  
[root@ip-172-31-3-152 /]#
```

Your Nginx config

```
user nginx;  
worker_processes auto;  
error_log /var/log/nginx/error.log;  
pid /run/nginx.pid;  
  
# Load dynamic modules. See /usr/share/doc/nginx/README.dynamic.  
include /usr/share/nginx/modules/*.conf;  
  
events {  
    worker_connections 1024;  
}  
  
http {  
    log_format main '$remote_addr - $remote_user [$time_local] "$request" ' '  
        '$status $body_bytes_sent "$http_referer" ' '  
        '"$http_user_agent" "$http_x_forwarded_for"';  
  
    access_log /var/log/nginx/access.log main;  
  
    sendfile        on;  
    tcp_nopush      on;  
    tcp_nodelay     on;  
    keepalive_timeout 65;  
    types_hash_max_size 4096;  
  
    include /etc/nginx/mime.types;  
    default_type application/octet-stream;  
  
    # Load modular configuration files from the /etc/nginx/conf.d directory.  
    # See http://nginx.org/en/docs/nginx_core_module.html#include  
    # for more information.  
    include /etc/nginx/conf.d/*.conf;  
}  
  
server {  
    listen 80;  
    server_name 52.66.235.166;  
  
    location / {  
        proxy_pass http://127.0.0.1:8000;  
        proxy_set_header Host $host;  
        proxy_set_header X-Real-IP $remote_addr;  
    }  
}
```

Browser output of the endpoint



2. What security measures would you take to harden an EC2 instance serving financial APIs?

1) Instance-Level Security

Use a Hardened OS, Restrict SSH Access

2) Network Security

Security Groups , VPC Configuration

3) Application-Level Security

HTTPS with TLS , API Authentication & Authorization , Input Validation & Rate Limiting

4) Monitoring & Logging

Enable Logging

5) IAM & Access Control

- 6) Backups & Recovery - Amazon EBS snapshots

Section 2: AWS Infra

3. Explain how you would:

Configure a PostgreSQL database using RDS
Connect it securely to the EC2 instance
Back it up regularly

Configure a PostgreSQL Database Using RDS

- 1) Create Database by using the **PostgreSQL Database -Standard**
- 2) Choose the **Available versions of pg sql**
- 3) Choose **Production** for high availability, or **Dev/Test** for development environments
- 4) Set a **DB Instance Name** Make the settings and the credential access settings
- 5) Choose DB Instance Class t3.micro
- 6) Choose the availability and durability
- 7) Connect to the instance and choose the vpc and private - subnets as per the instance
- 8) Not to grant for public access for (security)
- 9) Get Back ups on regular basis

Connect RDS PostgreSQL Securely to EC2

- 1) Security Group Configuration - inbound rule - port (5432)
- 2) Connect to Terminal

```
sudo yum install postgresql -y  
psql -h <endpoint> -U <username> -d <db-name> -W
```

- 3) Use Secrets Manager
- 4) Use SSL certificates from AWS RDS documentation

Back It Up Regularly

- 1) Set up backup retention period
- 2) AWS automatically performs daily snapshots

- 3) Ensures high availability by replicating data to a standby instance in a different AZ
- 4) Use snapshot copy to replicate backups across regions for disaster recovery
- 5) Use **AWS CloudWatch** to monitor backup operation
- 6) Set alarms for failed backups or storage thresholds

4. If your FastAPI service uploads files to S3:

How would you give it minimum permissions to upload to only a specific folder?

- 1) Create an IAM Role
- 2) Attach this role to your EC2 instance running FastAPI.
- 3) Limit Permissions :
 - a) allow s3:PutObject to the specific folder like uploads/
 - b) allow s3:ListBucket need to verify object existence

Share a sample IAM policy

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "AllowPutObjectOnlyToSpecificFolder",
      "Effect": "Allow",
      "Action": "s3:PutObject",
      "Resource": "arn:aws:s3:::mys3bucket1705/uploads/*"
    },
    {
      "Sid": "AllowListBucketForUploadsPrefix",
      "Effect": "Allow",
      "Action": "s3:ListBucket",
      "Resource": "arn:aws:s3:::mys3bucket1705",
      "Condition": {
        "StringLike": {
          "s3:prefix": "uploads/*"
        }
      }
    }
  ]
}
```

```

    }
  }
]
}

```

Boto3 for S3 access

```

import boto3

s3 = boto3.client('s3')
bucket_name = "mys3bucket1705"

def upload_file_to_s3(file, filename):
    s3.upload_fileobj(file, bucket_name, f"uploads/{filename}")

```

Section 3: Monitoring & Automation

**5. What tools would you use to:
Monitor uptime and errors?**

To monitor uptime and detect errors in a **production-grade FastAPI application on EC2**

- 1) Amazon CloudWatch
- 2) Prometheus + Grafana

Trigger alerts if memory or CPU usage goes beyond 80%?

To trigger alerts when system resources exceed a threshold -

- 1) Amazon CloudWatch Alarms

set EC2 **CPUUtilization** metri

Configure **threshold** at > 80% for CPU or memory

Trigger **SNS notification** to email, SMS, or Lambda.

- 2) CloudWatch Agent - Install on EC2 to collect **custom metrics** like memory, disk, and processes.

6. Create a systemd unit file that ensures the Gunicorn-based FastAPI app starts on reboot and restarts on failure

1) Create the Unit File

```
sudo vim /etc/systemd/system/fastapi.service
```

```
[Unit]
```

```
Description=Gunicorn instance to serve FastAPI app
```

```
After=network.target
```

```
[Service]
```

```
User=ec2-user
```

```
Group=ec2-user
```

```
WorkingDirectory=/home/ec2-user/fastapi-app
```

```
ExecStart=/home/ec2-user/fastapi-app/venv/bin/gunicorn main:app \
```

```
--workers 3 \
```

```
--bind 127.0.0.1:8000 \
```

```
-k uvicorn.workers.UvicornWorker
```

```
Restart=on-failure
```

```
RestartSec=5
```

```
[Install]
```

```
WantedBy=multi-user.target
```

```
sudo systemctl daemon-reexec
```

```
sudo systemctl daemon-reload
```

```
sudo systemctl enable fastapi.service
```

```
sudo systemctl start fastapi.service
```

```
sudo systemctl status fastapi.service
```

OUTPUT

```
Active: active (running)
```



```
[root@ip-172-31-3-152 system]# vi fastapi.service
[root@ip-172-31-3-152 system]# cat fastapi.service
[Unit]
Description=Gunicorn instance to serve FastAPI app
After=network.target

[Service]
User=ec2-user
Group=ec2-user
WorkingDirectory=/home/ec2-user/fastapi-app
ExecStart=/home/ec2-user/fastapi-app/venv/bin/gunicorn main:app \
    --workers 3 \
    --bind 127.0.0.1:8000 \
    -k uvicorn.workers.UvicornWorker
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target

[root@ip-172-31-3-152 system]# sudo systemctl daemon-reexec
[root@ip-172-31-3-152 system]# sudo systemctl daemon-reload
[root@ip-172-31-3-152 system]# sudo systemctl enable fastapi.service
Created symlink from /etc/systemd/system/multi-user.target.wants/fastapi.service to /etc/systemd/system/fastapi.service.
[root@ip-172-31-3-152 system]# sudo systemctl start fastapi.service
[root@ip-172-31-3-152 system]# sudo systemctl start fastapi.service
[root@ip-172-31-3-152 system]# sudo systemctl status fastapi.service
● fastapi.service - Gunicorn instance to serve FastAPI app
   Loaded: loaded (/etc/systemd/system/fastapi.service; enabled; vendor preset: disabled)
   Active: activating (auto-restart) (Result: exit-code) since Sat 2025-05-17 05:34:18 UTC; 3s ago
     Process: 5353 ExecStart=/home/ec2-user/fastapi-app/venv/bin/gunicorn main:app --workers 3 --bind 127.0.0.1:8000 -k uvicorn.workers.UvicornWorker (code=exited, status=200/CHDIR)
    Main PID: 5353 (code=exited, status=200/CHDIR)

May 17 05:34:18 ip-172-31-3-152.ap-south-1.compute.internal systemd[1]: fastapi.service: main process exited, code=exited, status=200/CHDIR
May 17 05:34:18 ip-172-31-3-152.ap-south-1.compute.internal systemd[1]: Unit fastapi.service entered failed state.
May 17 05:34:18 ip-172-31-3-152.ap-south-1.compute.internal systemd[1]: fastapi.service failed.
[root@ip-172-31-3-152 system]#
```

Section 4: Bonus

7. What are 2 cost optimization steps you'd recommend if this platform grows to handle 500+ MSMEs?

1. Use Auto Scaling + Spot Instances for EC2

Regular on-demand EC2 instances can be expensive at scale. So , Use **Auto Scaling Groups** to dynamically add/remove EC2 instances based on CPU or memory usage.

Mix in Spot Instances (up to 90% cheaper) for non-critical or stateless workloads like background workers or batch jobs.

2. Move to Serverless - (API + S3 + RDS)

. **Use API Gateway + Lambda** for request handling

. Store static content (files, images, docs) in S3 + CloudFront

. CloudFront speeds up content delivery to MSME users across

. Enable **RDS Auto Scaling**, turn on **storage autoscaling**